

Gas mentorship

Comprehensive guide about gas
optimization techniques

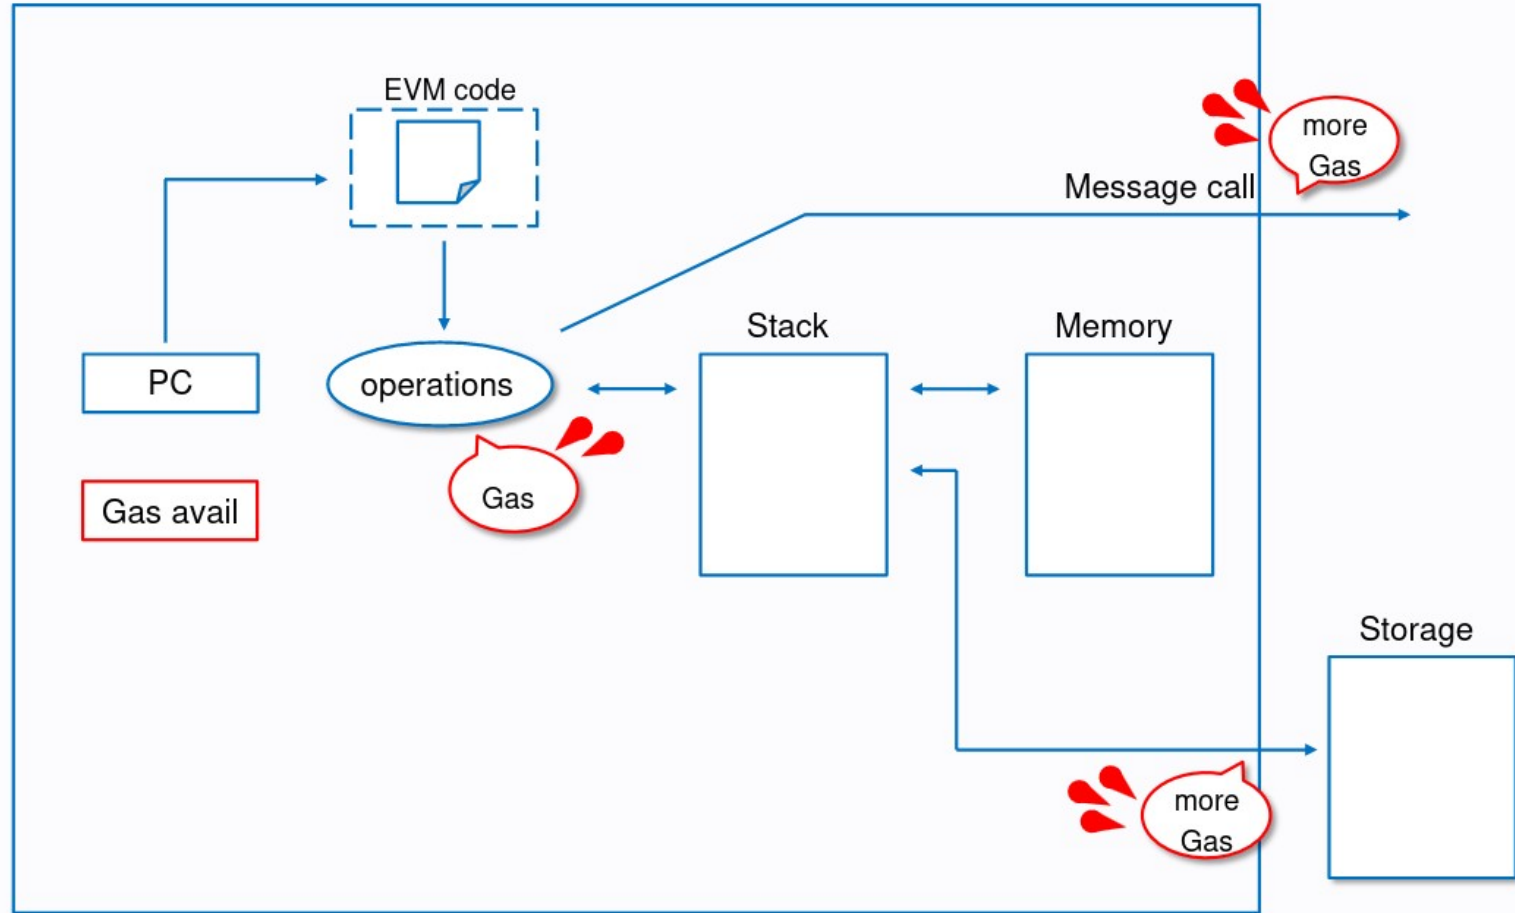
Content

- Introduction
 - General overview of gas in ETH
- EVM Basics
 - Background knowledge about the EVM
- Gas optimizations
 - In depth review of gas optimizations

Introduction

- What is gas?
 - Gas is the unit of computation in Ethereum
 - Each opcode has a fixed gas cost
 - Additionally, some opcodes have an dynamic gas cost on top of the fixed gas cost, e.g.
 - SLOAD/SSTORE
 - MLOAD/MSTORE

EVM



Introduction

- Each transaction has an intrinsic cost of 21k gas
- The more complex the transaction, the more expensive it is:
 - ETH Transfer: 21k gas
 - ERC20 Transfer: ~65k gas
 - ERC721 Transfer: ~84k gas
 - Uniswap V3 Swap: ~184k gas
 - Tornado Cash Deposit: ~1M gas

Introduction

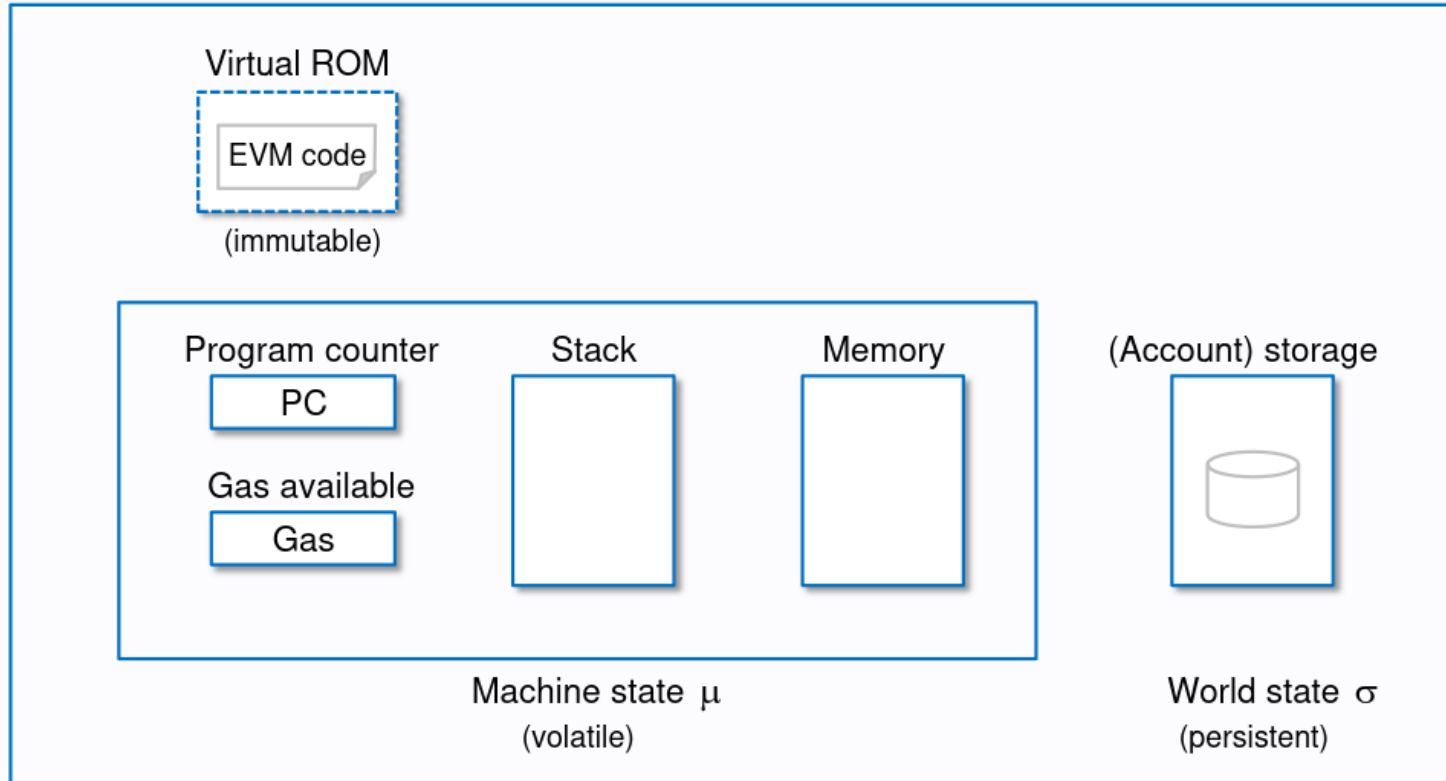
- How are gas costs calculated?
 - Since EIP-1559 the gas price consists of a
 - Base fee
 - Priority fee
 - The fee paid in USD is calculated as follows:

$(\text{gas used} \times \text{gas price (in gwei)} \times \text{ether price (in USD)}) / 1 \text{ billion}$

EVM Basics


- Stack based computer
- Deterministic, always the same output for the same input
- Creates an execution context when executing a smart contract
 - Contains several data regions
 - Program counter (PC)
 - Gas available
 - ...

Ethereum Virtual Machine (EVM)




Stack

- LIFO data structure
- PUSH (3 gas)
- POP (2 gas)



Memory

- Byte array
- MSTORE: 32 byte / 1 byte
- MLOAD: 32 byte



Memory opcodes

- Gas cost: 3 (static) + dynamic (memory expansion)

0x51	MLOAD	1	1	Load word from memory. $\mu'_s[0] \equiv \mu_m[\mu_s[0] \dots (\mu_s[0] + 31)]$ $\mu'_i \equiv \max(\mu_i, \lceil (\mu_s[0] + 32) \div 32 \rceil)$ The addition in the calculation of μ'_i is not subject to the 2^{256} modulo.
0x52	MSTORE	2	0	Save word to memory. $\mu'_m[\mu_s[0] \dots (\mu_s[0] + 31)] \equiv \mu_s[1]$ $\mu'_i \equiv \max(\mu_i, \lceil (\mu_s[0] + 32) \div 32 \rceil)$ The addition in the calculation of μ'_i is not subject to the 2^{256} modulo.
0x53	MSTORE8	2	0	Save byte to memory. $\mu'_m[\mu_s[0]] \equiv (\mu_s[1] \bmod 256)$ $\mu'_i \equiv \max(\mu_i, \lceil (\mu_s[0] + 1) \div 32 \rceil)$ The addition in the calculation of μ'_i is not subject to the 2^{256} modulo.

Memory expansion

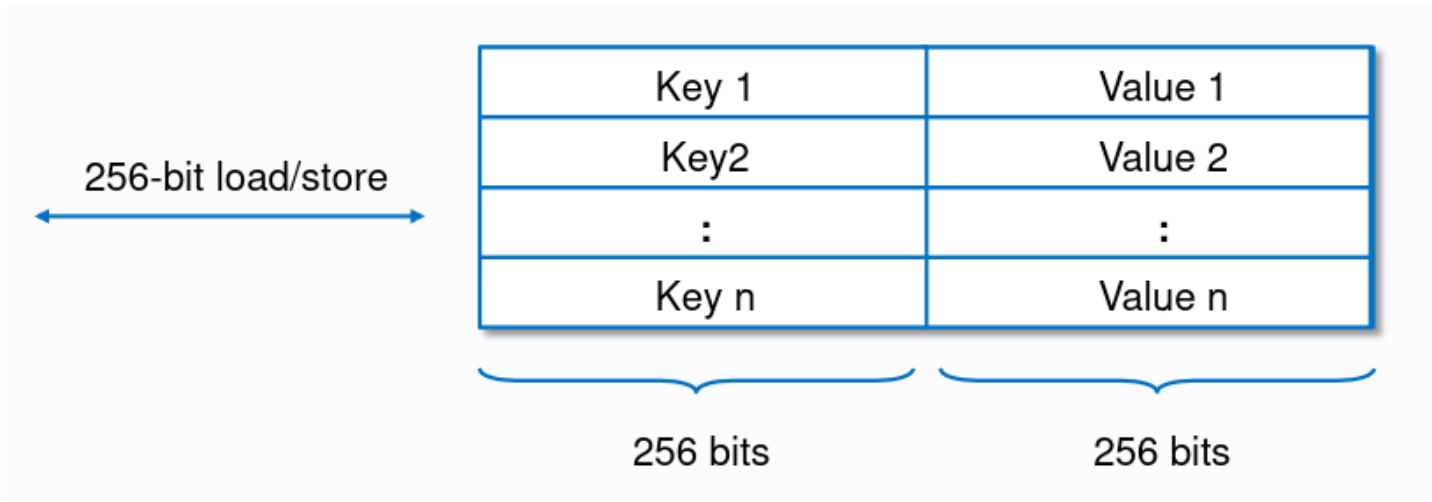
- When an offset is first accessed the memory is expanded in 32 byte chunks
- The cost of memory expansion grows quadratically:

$$C_{\text{mem}}(a) \equiv G_{\text{memory}} \cdot a + \left\lfloor \frac{a^2}{512} \right\rfloor$$

(`a` is the size of the memory, G_memory is 3)

Storage

- 256 bit key-value store, all values initialized as zero
- Multiple variables can be stored in a single slot (storage packing)



Storage: Gas

- Static cost of 100 gas for read (SLOAD) and write (SSTORE)
- Dynamic cost related to the state of the storage slot

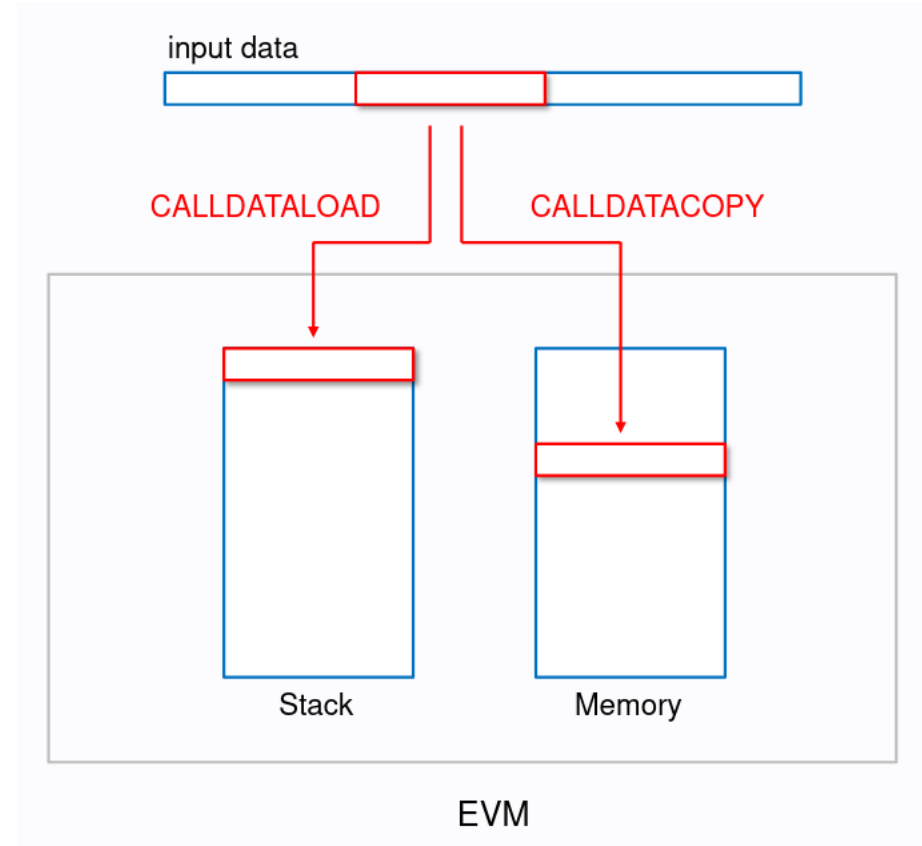
$G_{\text{warmaccess}}$	100	Cost of a warm account or storage access.
$G_{\text{accesslistaddress}}$	2400	Cost of warming up an account with the access list.
$G_{\text{accessliststorage}}$	1900	Cost of warming up a storage with the access list.
$G_{\text{coldaccountaccess}}$	2600	Cost of a cold account access.
$G_{\text{coldstorage}}$	2100	Cost of a cold storage access.
G_{sset}	20000	Paid for an SSTORE operation when the storage value is set to non-zero from zero.
G_{sreset}	2900	Paid for an SSTORE operation when the storage value's zeroness remains unchanged or is set to zero.
R_{sclear}	15000	Refund given (added into refund counter) when the storage value is set to zero from non-zero.

Storage: Gas

- Storage writes
 - zero to nonzero: 20k gas + 2.1k gas if first time access
 - Nonzero to nonzero: 5k gas
 - Nonzero to zero: refund up to 4.8k gas per storage variable
 - Zero to zero: 2.2k gas
- Storage reads
 - Cold: 2.1k gas
 - Warm: 100 gas

Calldata

- Similar to memory, but read only
- Can only be used in external function calls



Calldata: Gas

- Gas cost: 3 for calldataload/copy, 2 for calldatasize
- A byte of calldata costs 4 gas (if its zero) or 16 gas (nonzero)

0x35	CALLDATALOAD	1	1	Get input data of current environment. $\mu'_s[0] \equiv I_d[\mu_s[0] \dots (\mu_s[0] + 31)]$ with $I_d[x] = 0$ if $x \geq \ I_d\ $ This pertains to the input data passed with the message call instruction or transaction.
0x36	CALLDATASIZE	0	1	Get size of input data in current environment. $\mu'_s[0] \equiv \ I_d\ $ This pertains to the input data passed with the message call instruction or transaction.
0x37	CALLDATACOPY	3	0	Copy input data in current environment to memory. $\forall i \in \{0 \dots \mu_s[2] - 1\} : \mu'_m[\mu_s[0] + i] \equiv \begin{cases} I_d[\mu_s[1] + i] & \text{if } \mu_s[1] + i < \ I_d\ \\ 0 & \text{otherwise} \end{cases}$ The additions in $\mu_s[1] + i$ are not subject to the 2^{256} modulo. $\mu'_i \equiv M(\mu_i, \mu_s[0], \mu_s[2])$ This pertains to the input data passed with the message call instruction or transaction.

References

- [Ethereum EVM illustrated](#)
- [evm.codes](#)
- [Ethereum Yellowpaper](#)
- [Solidity internals](#)